

Calcul haute performance pour les systèmes embarqués (HPEC) Lab 3 - SIMD instruction sets

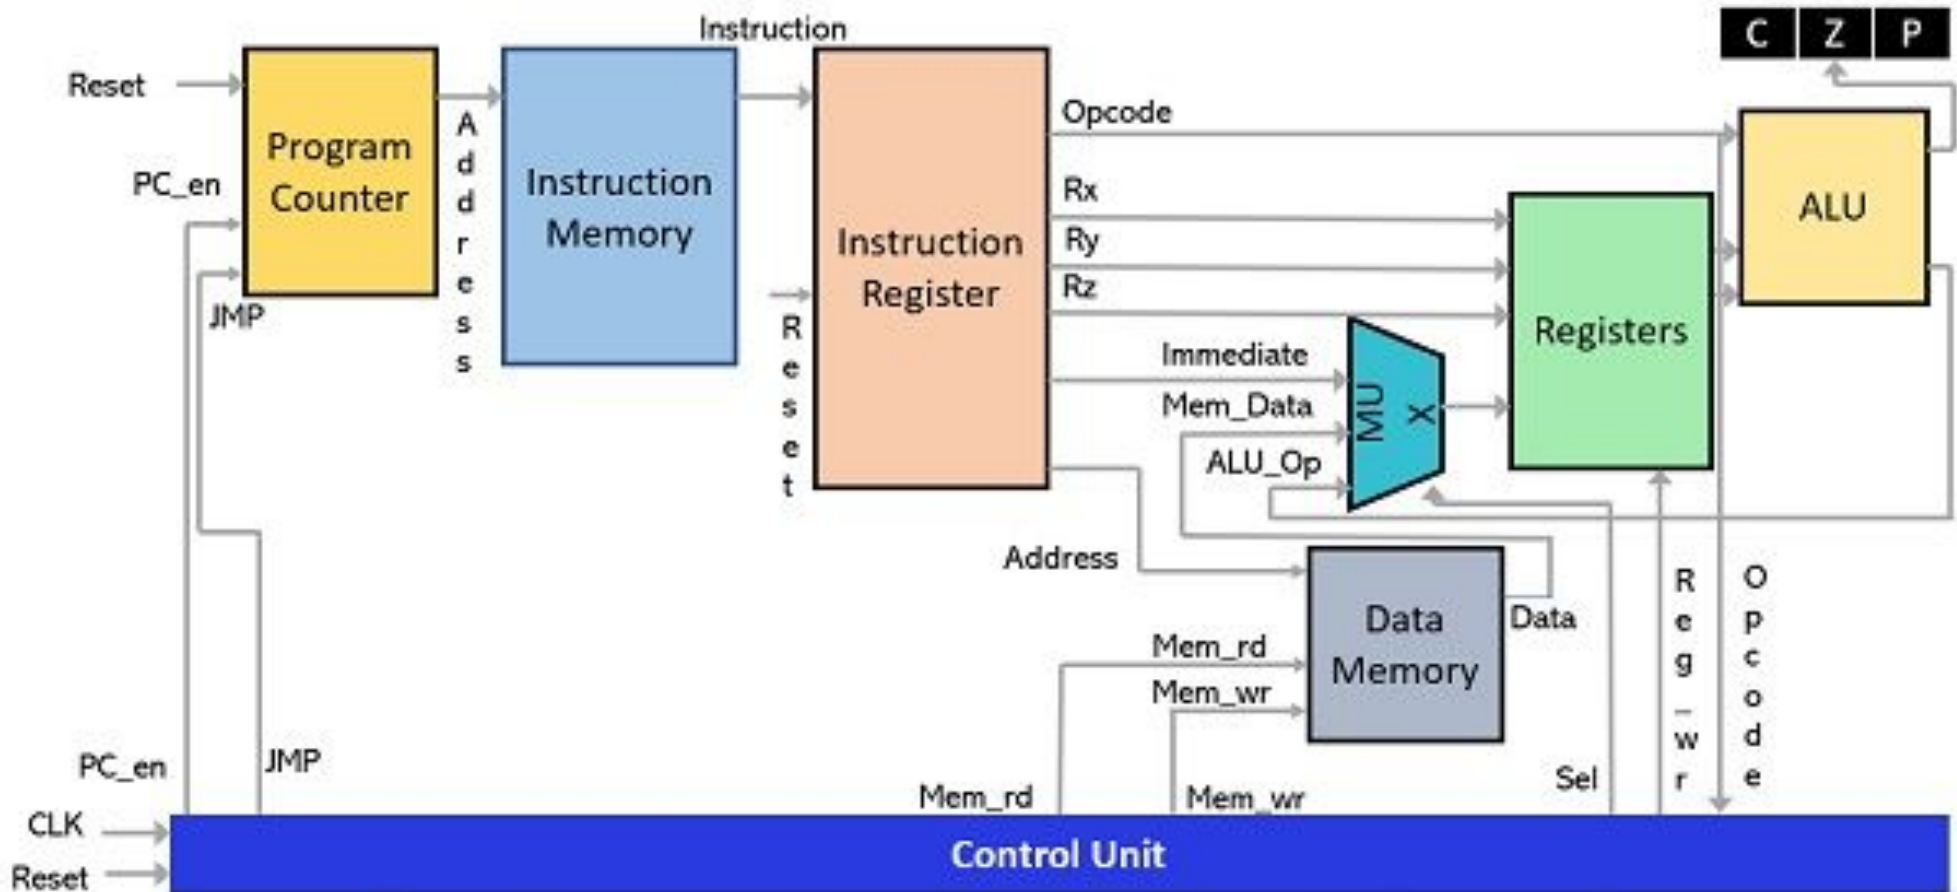
Lab 3 - SIMD instruction sets

Bertrand LE GAL

IMS laboratory, CNRS UMR 5218
Digital Circuits and Systems team
Bordeaux-INP, University of Bordeaux
France



Processor architecture from 90's



Increasing the hardware usage rate

- Current processors are:

- 64 bits word length,
- ≥ 18 pipeline stages,


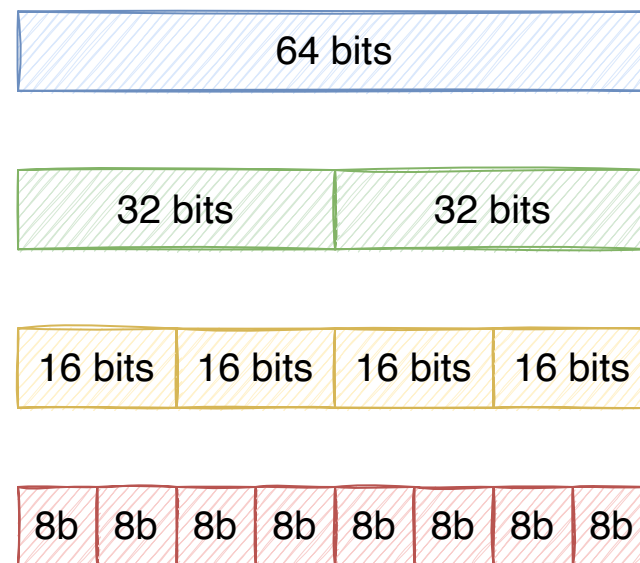
- It means that

- Registers are 64 bits wide,
- ALU resources are 64 bits wide,

- In real life

- Most data are ≤ 64 bits wide (pixels = 8b),
- Reusing hardware should be possible,


- It is easy for logical operation, more complex for others.



SIMD feature for parallel computing

Single Instruction Multiple Data (SIMD) paradigm

- Appears at the beginning of 20's,
- Needed for video processing,

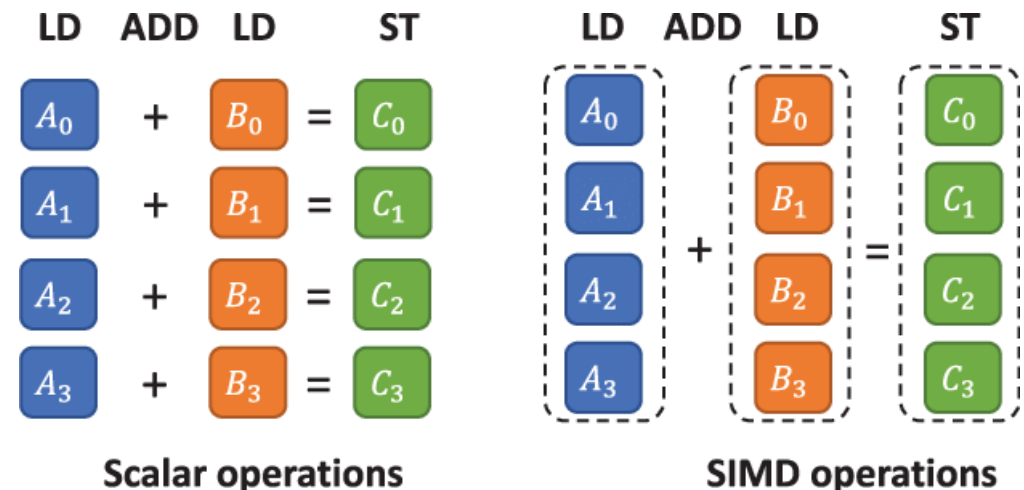


INTEL MMX ALU

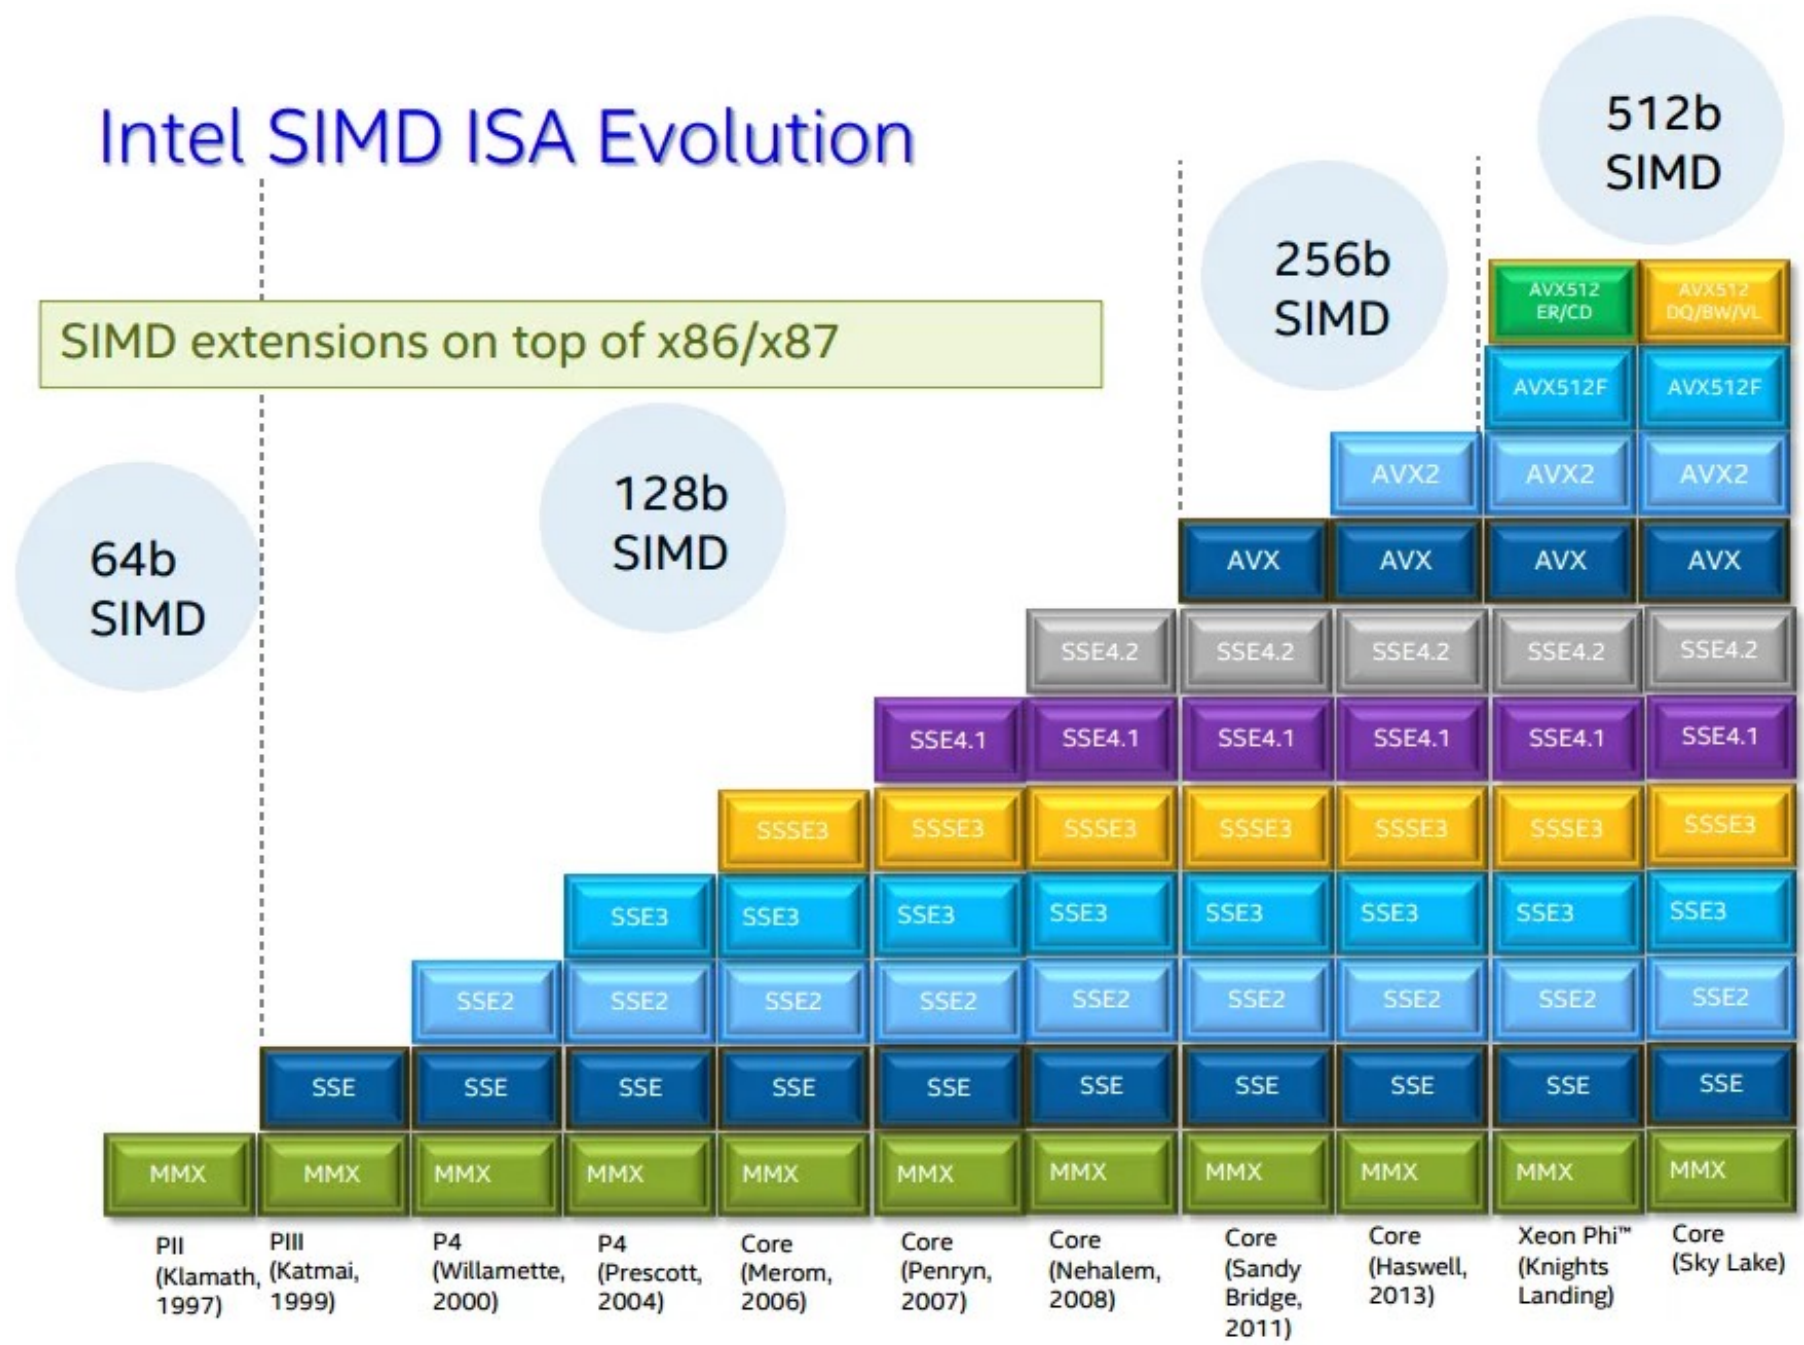
- 32 bits processor
- 64 bits integer ALU
- Dedicated register bank,

ALU operations

- Data width $\in \{8, 16, 32, 64\}$ bits.



Evolution des unités SIMD chez INTEL



Toy SIMD example - ARM Neon version

```
#include <arm_neon.h>

int vec_sum(
    float* dst,
    const float* src1,
    const float* src2,
    const int length)
{
    for (int i = 0; i < length; i += simd)
    {
        float v1 = src1[ i ]; // load (A) value
        float v2 = src2[ i ]; // load (B) value
        float re = v1 + v2;    // compute (A + B)
        dst[ i ] = re;         // store (C) value
    }
};
```

Toy SIMD example - ARM Neon version

```
#include <arm_neon.h>

int vec_sum(
    float* dst,
    const float* src1,
    const float* src2,
    const int length)
{
    int simd = sizeof(float32x4_t) / sizeof(float);

    for (int i = 0; i < length; i += simd)
    {
        float32x4_t v1 = vld1q_f32(src1 + i); // load (A) value
        float32x4_t v2 = vld1q_f32(src2 + i); // load (B) value
        float32x4_t re = vaddq_f32(v1, v2); // compute (A + B)
        vst1q_f32(dst + i, re); // store (C) value
    }
};
```

Toy SIMD example - INTEL SSE version

```
#include <immintrin.h>

int vec_sum(
    float* dst,
    const float* src1,
    const float* src2,
    const int length)
{
    int simd = sizeof(__m128) / sizeof(float); // = 4

    for (int i = 0; i < length; i += simd)
    {
        __m128 v1 = _mm_loadu_ps(src1 + i); // load (A) value
        __m128 v2 = _mm_loadu_ps(src2 + i); // load (B) value
        __m128 re = _mm_add_ps(v1, v2); // compute (A + B)
        _mm_storeu_ps(dst + i, re); // store (C) value
    }
};
```


Toy SIMD example - INTEL AVX version

```
#include <immintrin.h>

int vec_sum(
    float* dst,
    const float* src1,
    const float* src2,
    const int length)
{
    int simd = sizeof(__m256) / sizeof(float); // = 8

    for (int i = 0; i < length; i += simd)
    {
        __m256 v1 = _mm256_loadu_ps(src1 + i); // load (A) value
        __m256 v2 = _mm256_loadu_ps(src2 + i); // load (B) value
        __m256 re = _mm256_add_ps(v1, v2); // compute (A + B)
        _mm256_storeu_ps(dst + i, re); // store (C) value
    }
};
```

SIMD data types available in INTEL/ARM

ARM Neon

```
float32x4_t : 4x 32b float values (float)
float64x2_t : 2x 64b float values (double)

int64x2_t : 2x 64b integer values (int)
int32x4_t : 4x 32b integer values (int)
int16x8_t : 8x 16b integer values (int)
int8x16_t : 16x 8b integer values (int)

uint64x2_t : 2x 64b integer values (uint)
..... : .. ... .....
uint8x16_t : 16x 8b integer values (uint)
```

INTEL SIMD

```
__m128 : 4x 32b float values (float)
__m128d : 2x 64b float values (double)
__m128i : 2x 64b integer values (int/uint)
      : 4x 32b integer values (int/uint)
      : 8x 16b integer values (int/uint)
      : 16x 8b integer values (int/uint)

__m256 : 4x 32b float values (float)
__m256d : 2x 64b float values (double)
__m256i : 4x 64b integer values (int/uint)
      : 8x 32b integer values (int/uint)
      : 16x 16b integer values (int/uint)
      : 32x 8b integer values (int/uint)
```

The INTEL intrinsics for SSE/AVX targets

```
#include <xmmintrin.h>

// Function prefix depends on executed instruction

__m128 v1 = _mm_load_ps (mem_addr);    // f32 - load => WARNING ALIGNED !
__m128 v2 = _mm_add_ps  (v1, v1);      // f32 - add
__m128 v3 = _mm_mul_ps  (v1, v1);      // f32 - mul
__m128 v4 = _mm_fmadd_ps(v1, v2, v3);  // f32 - acc = v3 + v1 * v2
__m128 v5 = _mm_max_ps  (v1, v4);      // f32 - max
_mm_store_ps (mem_addr, v5);           // f32 - store

// Function suffix depends on processed data

__m128  _mm_add_ps      (__m128 a, __m128 b); // 128b
__m128d _mm_add_pd      (__m128d a, __m128d b); // 128b
__m128i _mm_add_epi8    (__m128i a, __m128i b); // 128b
__m128i _mm_add_epi16   (__m128i a, __m128i b); // 128b
__m128i _mm_add_epi32   (__m128i a, __m128i b); // 128b
__m128i _mm_add_epi64   (__m128i a, __m128i b); // 128b

__m256  _mm256_add_ps   (__m256 a, __m256 b); // 256b
// ... ..

__m512  _mm512_add_ps   (__m512 a, __m512 b); // 512b
// ... ..
```

<https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>

The ARM intrinsics for NEON targets

```
#include <arm_neon.h>

// Function prefix depends on executed instruction

float32x4_t v1    = vld1q_f32(values);      // f32 - load
float32x4_t sum   = vaddq_f32(v1, v2);      // f32 - add
float32x4_t prod  = vmulq_f32(v1, v2);      // f32 - mul
float32x4_t acc   = vmlaq_f32(v3, v1, v2);  // f32 - acc = v3 + v1 * v2
float32x4_t v2    = vmaxq_f32(v0, v1);      // f32 - max
vst1q_f32(values, v1);                      // f32 - store

// Function suffix depends on processed data

int8x8_t      vpadd_s8 (int8x8_t a,      int8x8_t  b);
int16x4_t     vpadd_s16(int16x4_t a,     int16x4_t  b);
int32x2_t     vpadd_s32(int32x2_t a,     int32x2_t  b);
uint8x8_t     vpadd_u8 (uint8x8_t a,     uint8x8_t  b);
uint16x4_t    vpadd_u16(uint16x4_t a,    uint16x4_t  b);
uint32x2_t    vpadd_u32(uint32x2_t a,    uint32x2_t  b);
float32x2_t   vpadd_f32(float32x2_t a,   float32x2_t b);
```

<https://github.com/thenifty/neon-guide>

<https://developer.arm.com/documentation/dui0472/m/Using-NEON-Support>

Missing features in INTEL ISA

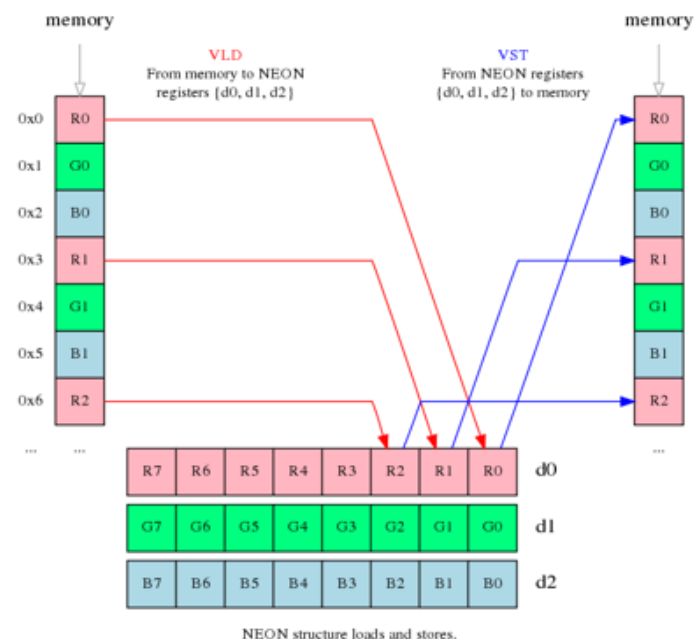

● INTEL processors

- aligned or unaligned two instr.
- contiguous memory load/store

● ARM processors

- aligned/unaligned one instr.
- contiguous memory load/store
- interleaved load/store

● On intel architecture such feature should be hand implemented



<https://community.arm.com/arm-community-blogs/b/architectures-and-processors-blog/posts/coding-for-neon---part-i-load-and-stores>

Missing features in ARM NEON

Complex algorithms & memory structures needs at runtime


- Data permutation,
- Data duplication, etc,
- Conditional moves.

Two intrinsic families

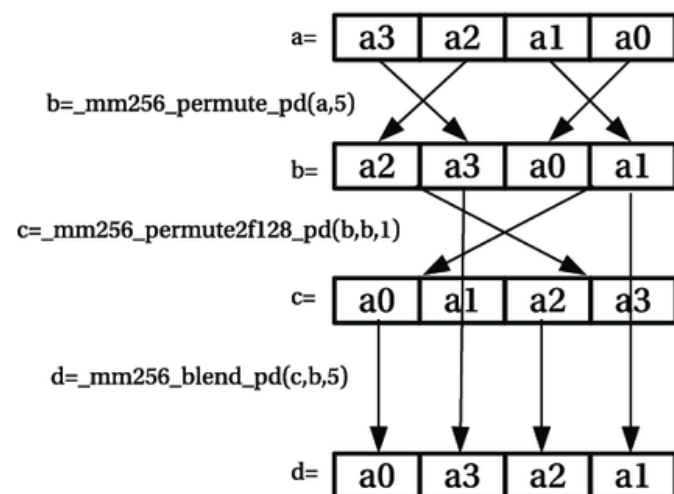
- `_mm*_shuffle_ps`
- `_mm*_permute_ps`
- Easy to use with one lane, otherwise...

```
__m256i _mm256_blendv_epi8 (__m256i a,  
                             __m256i b,  
                             __m256i mask)
```

```
res = _mm256_permute_ps(vec, 0b01110100)
```



Cyclic double vector rotate




How to compile codes with SIMD intrinsics

⦿ Depends on the targeted CPU or toolchain

- `g++ -o main main.cpp -O3 -mtune=native -march=native` # INTEL CPUs
- `g++ -o main main.cpp -O3 -mtune=native -mcpu=native` # ARM CPUs

⦿ The « native » architecture defines some preprocessor directives

- `__SSE4_2__ __AVX__ __AVX2__ __AVX512F__` # INTEL CPUs
- `__ARM_NEON__ __ARM_NEON` # ARM CPUs



Compiler auto-vectorization features


• Compilers can do the job 4u

• Pro

- **Generic** (int8_t ... int32_t ... float),
- **Generic** (SSE4 ... AVX2 ... AVX512),
- **Generic** (ARM & INTEL),
- (Very) **efficient** thanks to cost models used by the compiler,
- Compiler provides vectorization tips.

• Con

- **Limited** to simple kernels,
- **Learn compiler code style**,
- Needs additional pragma or keywords.



<https://gcc.gnu.org/projects/tree-ssa/vectorization.html>

<https://llvm.org/docs/Vectorizers.html>

<https://blog.minhazav.dev/guide-compiler-to-auto-vectorise/>

Managing « efficiently » various SIMD ISA

Writing SIMD codes

- long and painful,
- adapted to target features,

C++ wrappers for SIMD intrinsics

- **XSIMD**
- **VectorClass**
- **MIPPS**
- **nova-simd** and ...

Easy software codes

- Portability & flexibility
- What is really executed ?

```
#include <stddef>
#include <vector>
#include "xsimd/xsimd.hpp"

namespace xs = xsimd;
using vector_type = std::vector<double, xsimd::aligned_allocator<double>>;

void mean(const vector_type& a, const vector_type& b, vector_type& res)
{
    std::size_t size = a.size();
    constexpr std::size_t simd_size = xsimd::simd_type<double>::size;
    std::size_t vec_size = size - size % simd_size;

    for(std::size_t i = 0; i < vec_size; i += simd_size)
    {
        auto ba = xs::load_aligned(&a[i]);
        auto bb = xs::load_aligned(&b[i]);
        auto bres = (ba + bb) / 2;
        bres.store_aligned(&res[i]);
    }
    for(std::size_t i = vec_size; i < size; ++i)
    {
        res[i] = (a[i] + b[i]) / 2;
    }
}
```